

Wine_Quality_Prediction

April 30, 2023

1 WINE QUALITY PREDICTION

Importing The Libraries

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score
```

Read the csv file into a DataFrame

```
[3]: df=pd.read_csv('wine.csv')
```

```
[4]: df.head(10)
```

```
[4]:  fixed acidity  volatile acidity  citric acid  residual sugar  chlorides  \
0           7.4             0.70         0.00           1.9       0.076
1           7.8             0.88         0.00           2.6       0.098
2           7.8             0.76         0.04           2.3       0.092
3          11.2             0.28         0.56           1.9       0.075
4           7.4             0.70         0.00           1.9       0.076
5           7.4             0.66         0.00           1.8       0.075
6           7.9             0.60         0.06           1.6       0.069
7           7.3             0.65         0.00           1.2       0.065
8           7.8             0.58         0.02           2.0       0.073
9           7.5             0.50         0.36           6.1       0.071
```

```
    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates  \
0              11.0             34.0    0.9978  3.51       0.56
1              25.0             67.0    0.9968  3.20       0.68
2              15.0             54.0    0.9970  3.26       0.65
3              17.0             60.0    0.9980  3.16       0.58
4              11.0             34.0    0.9978  3.51       0.56
5              13.0             40.0    0.9978  3.51       0.56
```

6	15.0	59.0	0.9964	3.30	0.46
7	15.0	21.0	0.9946	3.39	0.47
8	9.0	18.0	0.9968	3.36	0.57
9	17.0	102.0	0.9978	3.35	0.80

	alcohol	quality
0	9.4	5
1	9.8	5
2	9.8	5
3	9.8	6
4	9.4	5
5	9.4	5
6	9.4	5
7	10.0	7
8	9.5	7
9	10.5	5

```
[5]: df.shape
```

```
[5]: (1599, 12)
```

```
[6]: df.isnull().sum()
```

```
[6]: fixed acidity      0
      volatile acidity  0
      citric acid      0
      residual sugar   0
      chlorides        0
      free sulfur dioxide 0
      total sulfur dioxide 0
      density          0
      pH              0
      sulphates        0
      alcohol          0
      quality          0
      dtype: int64
```

```
[8]: df.describe()
```

```
[8]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	8.319637	0.527821	0.270976	2.538806
std	1.741096	0.179060	0.194801	1.409928
min	4.600000	0.120000	0.000000	0.900000
25%	7.100000	0.390000	0.090000	1.900000
50%	7.900000	0.520000	0.260000	2.200000
75%	9.200000	0.640000	0.420000	2.600000

max	15.900000	1.580000	1.000000	15.500000
-----	-----------	----------	----------	-----------

	chlorides	free sulfur dioxide	total sulfur dioxide	density \
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	0.087467	15.874922	46.467792	0.996747
std	0.047065	10.460157	32.895324	0.001887
min	0.012000	1.000000	6.000000	0.990070
25%	0.070000	7.000000	22.000000	0.995600
50%	0.079000	14.000000	38.000000	0.996750
75%	0.090000	21.000000	62.000000	0.997835
max	0.611000	72.000000	289.000000	1.003690

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[9]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1599 entries, 0 to 1598
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          1599 non-null   float64
1   volatile acidity       1599 non-null   float64
2   citric acid            1599 non-null   float64
3   residual sugar         1599 non-null   float64
4   chlorides              1599 non-null   float64
5   free sulfur dioxide    1599 non-null   float64
6   total sulfur dioxide   1599 non-null   float64
7   density                1599 non-null   float64
8   pH                    1599 non-null   float64
9   sulphates              1599 non-null   float64
10  alcohol                1599 non-null   float64
11  quality                1599 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 150.0 KB
```

```
[10]: df.dtypes
```


```
[10]: fixed acidity      float64
      volatile acidity  float64
      citric acid       float64
      residual sugar    float64
      chlorides         float64
      free sulfur dioxide float64
      total sulfur dioxide float64
      density          float64
      pH              float64
      sulphates       float64
      alcohol         float64
      quality         int64
      dtype: object
```

Visualization of Data

Number of values for each quality

```
[16]: sns.catplot(x='quality', data = df, kind = 'count')
```


```
[16]: <seaborn.axisgrid.FacetGrid at 0x7feb6a606a70>
```



Volatile Acidity vs Quality

```
[18]: plot = plt.figure(figsize=(10,5))  
sns.barplot(x='quality', y = 'volatile acidity', data = df)
```


```
[18]: <Axes: xlabel='quality', ylabel='volatile acidity'>
```



Citric Acid vs Quality

```
[19]: plot = plt.figure(figsize=(10,5))  
sns.barplot(x='quality', y = 'citric acid', data = df)
```


```
[19]: <Axes: xlabel='quality', ylabel='citric acid'>
```



Residual Sugar vs Quality

```
[20]: plot = plt.figure(figsize = (10,5))  
sns.barplot(x = 'quality', y = 'residual sugar', data = df)
```


```
[20]: <Axes: xlabel='quality', ylabel='residual sugar'>
```



Chlorides vs Quality

```
[22]: fig = plt.figure(figsize = (10,5))
sns.barplot(x = 'quality', y = 'chlorides', data = df)
#Composition of chloride also go down as we go higher in the quality of the wine
```


```
[22]: <Axes: xlabel='quality', ylabel='chlorides'>
```



Free Sulphur Dioxide vs Quality

```
[23]: fig = plt.figure(figsize = (10,5))
sns.barplot(x = 'quality', y = 'free sulfur dioxide', data = df)
```

```
[23]: <Axes: xlabel='quality', ylabel='free sulfur dioxide'>
```




Correlation

```
[24]: Corr = df.corr()
```

```
[25]: plt.figure(figsize=(15,10))
sns.heatmap(Corr, cbar=True, square=True, fmt = '.1f', annot = True,
            annot_kws={'size':8})
```

```
[25]: <Axes: >
```

Data Preprocessing

Separate the Data and label

```
[26]: X = df.drop('quality',axis=1)
```

```
[27]: print(X)
```

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	\
0	7.4		0.700		1.9	0.076
1	7.8		0.880		2.6	0.098
2	7.8		0.760		2.3	0.092
3	11.2		0.280		1.9	0.075
4	7.4		0.700		1.9	0.076
...	
1594	6.2		0.600		2.0	0.090
1595	5.9		0.550		2.2	0.062
1596	6.3		0.510		2.3	0.076

1597	5.9	0.645	0.12	2.0	0.075
1598	6.0	0.310	0.47	3.6	0.067

	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates \
0	11.0	34.0	0.99780	3.51	0.56
1	25.0	67.0	0.99680	3.20	0.68
2	15.0	54.0	0.99700	3.26	0.65
3	17.0	60.0	0.99800	3.16	0.58
4	11.0	34.0	0.99780	3.51	0.56
...
1594	32.0	44.0	0.99490	3.45	0.58
1595	39.0	51.0	0.99512	3.52	0.76
1596	29.0	40.0	0.99574	3.42	0.75
1597	32.0	44.0	0.99547	3.57	0.71
1598	18.0	42.0	0.99549	3.39	0.66

	alcohol
0	9.4
1	9.8
2	9.8
3	9.8
4	9.4
...	...
1594	10.5
1595	11.2
1596	11.0
1597	10.2
1598	11.0

[1599 rows x 11 columns]

```
[28]: Y = df['quality'].apply(lambda y_value: 1 if y_value>=7 else 0)
```

```
[29]: print(Y)
```

0	0
1	0
2	0
3	0
4	0
..	
1594	0
1595	0
1596	0
1597	0
1598	0

Name: quality, Length: 1599, dtype: int64

Train Test Split

```
[30]: X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2,
↳ random_state=3)
```

Model Training

```
[31]: model = RandomForestClassifier()
```

```
[32]: model.fit(X_train, Y_train)
```

```
[32]: RandomForestClassifier()
```

Evaluating The Model

Accuracy on test data

```
[33]: X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
[34]: print('Accuracy : ', test_data_accuracy)
```

Accuracy : 0.928125

```
[35]: input_data = (7.5,0.5,0.36,6.1,0.071,17.0,102.0,0.9978,3.35,0.8,10.5)
```

Changing the input data to a numpy array

```
[38]: input_data_to_numpy_array = np.asarray(input_data)
```

```
[39]: input_data_reshape = input_data_to_numpy_array.reshape(1,-1)
#Here we are reshaping the data as we are predicting the label for only one
↳ instance
```

Prediction

```
[40]: prediction = model.predict(input_data_reshape)
print(prediction)

if (prediction[0]==1):
    print('Good Quality Wine')
else:
    print('Bad Quality Wine')
```

[0]

Bad Quality Wine

/usr/local/lib/python3.10/dist-packages/sklearn/base.py:439: UserWarning: X does not have valid feature names, but RandomForestClassifier was fitted with feature

```
names
```

```
  warnings.warn(
```

```
[ ]:
```